

Problem and fix: the page title never reached the model

Hebrew RAG over Kol-Zchut · retrieval stage only · 24,487 paragraphs from 7,007 pages



Observation: "דרכון דחוף" (urgent passport)

- 1 ביומטרי (biometric) הוצאת דרכון
- 2 זמני (temporary): הוצאת דרכון correct

The two bodies are nearly identical. The difference is **only in the title**, and the model embeds only the paragraph body. Of 187 misses, 63 were at rank 2–3, mostly behind a sibling page.

The fix: the title becomes part of the embedded text

Before

נכים עשויים לקבל מהרשות המקומית הנחה בתשלומי הארנונה...

After

הנחה בארנונה לנכים ← the added title
נכים עשויים לקבל מהרשות המקומית הנחה בתשלומי הארנונה...

One config key. No retraining; search and API unchanged; removing the key restores the original.

Starting point, 296 questions

58%


right page among the 3 sent to GPT (hit@3)

42%

of questions: GPT gets 3 pages without the right one

Results: what ships, plus an optional LLM second stage

hit@3 = right page among the 3 sent to GPT · Webiks: your validation split · Clean: 150 questions on pages outside training, written with the title hidden



Version	Webiks: hit@3 / MRR	Clean: hit@3 / MRR	Added latency	Cost per 1,000 questions
Original system	58.1 / 0.496	74.7 / 0.653		
Title in text (shipped)	64.5 / 0.577	78.7 / 0.711	0	0
+ gpt-4o-mini, pick from 30 titles	75.3 / 0.692	84.0 / 0.757	0.73 s	~\$0.12
+ Jev, pick from 30 titles	78.0 / 0.710	82.0 / 0.719	0.44 s	\$0.07
+ Jev, title + short snippet	77.0 / 0.696	83.3 / 0.771	0.53 s	\$0.25

Verdict: title in text is what ships, and it improves both sets on both metrics. The second stage, off by default, adds more. Jev and gpt-4o-mini are **on par**: every 95% CI of the difference between them includes zero. Jev is faster and cheaper.

